

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [3]: from sklearn import preprocessing
from sklearn.model_selection import train_test_split
```

```
In [5]: from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score
```

```
In [9]: data=pd.read_csv(r'C:\Users\DELL\Downloads\car-mpg.csv')
```

```
In [11]: data.head()
```

```
Out[11]:
```

	mpg	cyl	displacement	horsepower	weight	acceleration	year	origin	car_type	car_name
0	18.0	8	307.0	130	3504	12.0	70	1	0	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	0	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	0	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	0	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	0	ford torino

```
In [13]: data = data.drop(['car_name'], axis=1)
```

```
In [15]: data['origin'] = data['origin'].replace({1: 'america', 2: 'europe', 3: 'asia'})
data= pd.get_dummies(data, columns =['origin'], dtype=int)
data= data.replace('?', np.nan)
```


```
In [19]: data=data.apply(pd.to_numeric, errors='ignore')
numeric_cols=data.select_dtypes(include=[np.number]).columns
data[numeric_cols]= data[numeric_cols].apply(lambda x: x.fillna(x.median()))
```

C:\Users\DELL\AppData\Local\Temp\ipykernel\_27080\202028991.py:1: FutureWarning: errors='ignore' is deprecated and will raise in a future version. Use to\_numeric without passing `errors` and catch exceptions explicitly instead  
data=data.apply(pd.to\_numeric, errors='ignore')

```
In [21]: data.head()
```

Out[21]:

	mpg	cyl	displacement	horsepower	weight	acceleration	year	car_type	origin_america	origin_asia	origin_europe
0	18.0	8	307.0	130.0	3504	12.0	70	0	1	0	0
1	15.0	8	350.0	165.0	3693	11.5	70	0	1	0	0
2	18.0	8	318.0	150.0	3436	11.0	70	0	1	0	0
3	16.0	8	304.0	150.0	3433	12.0	70	0	1	0	0
4	17.0	8	302.0	140.0	3449	10.5	70	0	1	0	0



## Model Building

```
In [23]: X=data.drop(['mpg'],axis=1)
y=data[['mpg']]
```

```
In [31]: X_s=preprocessing.scale(X)
X_columns=X.columns
X_s=pd.DataFrame(X_s, columns= X_columns)

y_s=preprocessing.scale(y)
y_columns=y.columns
y_s=pd.DataFrame(y_s, columns= y_columns)
```

```
In [33]: X_train, X_test, y_train,y_test = train_test_split(X_s, y_s, test_size = 0.30, r
X_train.shape
```

Out[33]: (278, 10)

## Simple Linear Regression

```
In [38]: regression_model = LinearRegression()
regression_model.fit(X_train, y_train)

for idx, col_name in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col_name, regression_model.coef_

intercept = regression_model.intercept_[0]
print('The intercept is {}'.format(intercept))
```

```
The coefficient for cyl is 0.3210223856916108
The coefficient for disp is 0.3248343091848394
The coefficient for hp is -0.2291695005943759
The coefficient for wt is -0.7112101905072299
The coefficient for acc is 0.014713682764191435
The coefficient for yr is 0.3755811949510741
The coefficient for car_type is 0.38147694842331
The coefficient for origin_america is -0.0747224754758417
The coefficient for origin_asia is 0.04451525203567813
The coefficient for origin_europe is 0.04834854953945371
The intercept is 0.019284116103639715
```

## Regularized Ridge Regression

```
In [41]: ridge_model = Ridge(alpha = 0.3)
ridge_model.fit(X_train, y_train)
```

```
print('Ridge model coef: {}'.format(ridge_model.coef_))
```

```
Ridge model coef: [[ 0.31649043  0.31320707 -0.22876025 -0.70109447  0.01295851
 0.37447352
 0.37725608 -0.07423624  0.04441039  0.04784031]]
```

Regularized Lasso Regression

```
In [44]: lasso_model = Lasso(alpha = 0.1)
lasso_model.fit(X_train, y_train)

print('Lasso model coef: {}'.format(lasso_model.coef_))
```

```
Lasso model coef: [-0.          -0.          -0.01690287 -0.51890013  0.
 0.28138241
 0.1278489  -0.01642647  0.          0.          ]
```

Score Comparison

```
In [47]: #Simple Linear Model
print(regression_model.score(X_train, y_train))
print(regression_model.score(X_test, y_test))

print('*****')

#Ridge
print(ridge_model.score(X_train, y_train))
print(ridge_model.score(X_test, y_test))

print('*****')

#Lasso
print(lasso_model.score(X_train, y_train))
print(lasso_model.score(X_test, y_test))
```


```
0.8343770256960538
0.8513421387780067
*****
0.8343617931312617
0.8518882171608501
*****
0.7938010766228453
0.8375229615977084
```

Model Parameter Tuning

```
In [52]: data_train_test = pd.concat([X_train, y_train], axis =1)
data_train_test.head()
```

Out[52]:

	cyl	disp	hp	wt	acc	yr	car_type	origin_i
<b>350</b>	-0.856321	-0.849116	-1.081977	-0.893172	-0.242570	1.351199	0.941412	0
<b>59</b>	-0.856321	-0.925936	-1.317736	-0.847061	2.879909	-1.085858	0.941412	-1
<b>120</b>	-0.856321	-0.695475	0.201600	-0.121101	-0.024722	-0.815074	0.941412	-1
<b>12</b>	1.498191	1.983643	1.197027	0.934732	-2.203196	-1.627426	-1.062235	0
<b>349</b>	-0.856321	-0.983552	-0.951000	-1.165111	0.156817	1.351199	0.941412	-1



In [54]: `import statsmodels.formula.api as smf`  
`ols1 = smf.ols(formula = 'mpg ~ cyl+disp+hp+wt+acc+yr+car_type+origin_america+or`  
`ols1.params`

Out[54]:

Intercept	0.019284
cyl	0.321022
disp	0.324834
hp	-0.229170
wt	-0.711210
acc	0.014714
yr	0.375581
car_type	0.381477
origin_america	-0.074722
origin_europe	0.048349
origin_asia	0.044515
dtype:	float64

In [56]: `print(ols1.summary())`

# OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.834
Model:	OLS	Adj. R-squared:	0.829
Method:	Least Squares	F-statistic:	150.0
Date:	Thu, 21 Aug 2025	Prob (F-statistic):	3.12e-99
Time:	13:52:20	Log-Likelihood:	-146.89
No. Observations:	278	AIC:	313.8
Df Residuals:	268	BIC:	350.1
Df Model:	9		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	0.0193	0.025	0.765	0.445	-0.030	0.069
cyl	0.3210	0.112	2.856	0.005	0.100	0.542
disp	0.3248	0.128	2.544	0.012	0.073	0.576
hp	-0.2292	0.079	-2.915	0.004	-0.384	-0.074
wt	-0.7112	0.088	-8.118	0.000	-0.884	-0.539
acc	0.0147	0.039	0.373	0.709	-0.063	0.092
yr	0.3756	0.029	13.088	0.000	0.319	0.432
car_type	0.3815	0.067	5.728	0.000	0.250	0.513
origin_america	-0.0747	0.020	-3.723	0.000	-0.114	-0.035
origin_europe	0.0483	0.021	2.270	0.024	0.006	0.090
origin_asia	0.0445	0.020	2.175	0.031	0.004	0.085

Omnibus:	22.678	Durbin-Watson:	2.105
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36.139
Skew:	0.513	Prob(JB):	1.42e-08
Kurtosis:	4.438	Cond. No.	1.04e+16

## Notes:



- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.44e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [58]: mse = np.mean((regression_model.predict(X_test)-y_test)**2)

import math
rmse = math.sqrt(mse)
print('Root Mean Squared Error: {}'.format(rmse))
```


Root Mean Squared Error: 0.3776693425408783


```
In [62]: fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['hp'], y= y_test['mpg'], color='green', lowess=True )
plt.show()
```



```
In [68]: fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['hp'], y= y_test['mpg'], color='green', lowess=True )
plt.show()


fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['acc'], y= y_test['mpg'], color='green', lowess=True )
plt.show()
```





```
In [72]: y_pred = regression_model.predict(X_test)

plt.scatter(y_test['mpg'], y_pred)
plt.show()
```



```
In [ ]:
```



