

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
In [3]: %matplotlib inline
```

```
In [7]: dataset= pd.read_csv(r'C:\Users\DELL\Downloads\20th - m1r\20th - m1r\MLR\House_d
dataset
```

```
Out[7]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7
2	5631500400	20150225T000000	180000.0	2	1.00	770	10
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8
...
21608	263000018	20140521T000000	360000.0	3	2.50	1530	1
21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5
21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1
21611	291310100	20150116T000000	400000.0	3	2.50	1600	2
21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1

21613 rows × 21 columns



```
In [9]: dataset.head()
```

```
Out[9]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080

5 rows × 21 columns



```
In [11]: print(dataset.isnull().any())
```

```

id                False
date              False
price             False
bedrooms          False
bathrooms         False
sqft_living       False
sqft_lot          False
floors            False
waterfront        False
view              False
condition         False
grade             False
sqft_above        False
sqft_basement     False
yr_built          False
yr_renovated      False
zipcode           False
lat               False
long              False
sqft_living15     False
sqft_lot15        False
dtype: bool

```

```
In [13]: print(dataset.dtypes)
```

```

id                int64
date              object
price             float64
bedrooms          int64
bathrooms         float64
sqft_living       int64
sqft_lot          int64
floors            float64
waterfront        int64
view              int64
condition         int64
grade             int64
sqft_above        int64
sqft_basement     int64
yr_built          int64
yr_renovated      int64
zipcode           int64
lat               float64
long              float64
sqft_living15     int64
sqft_lot15        int64
dtype: object

```

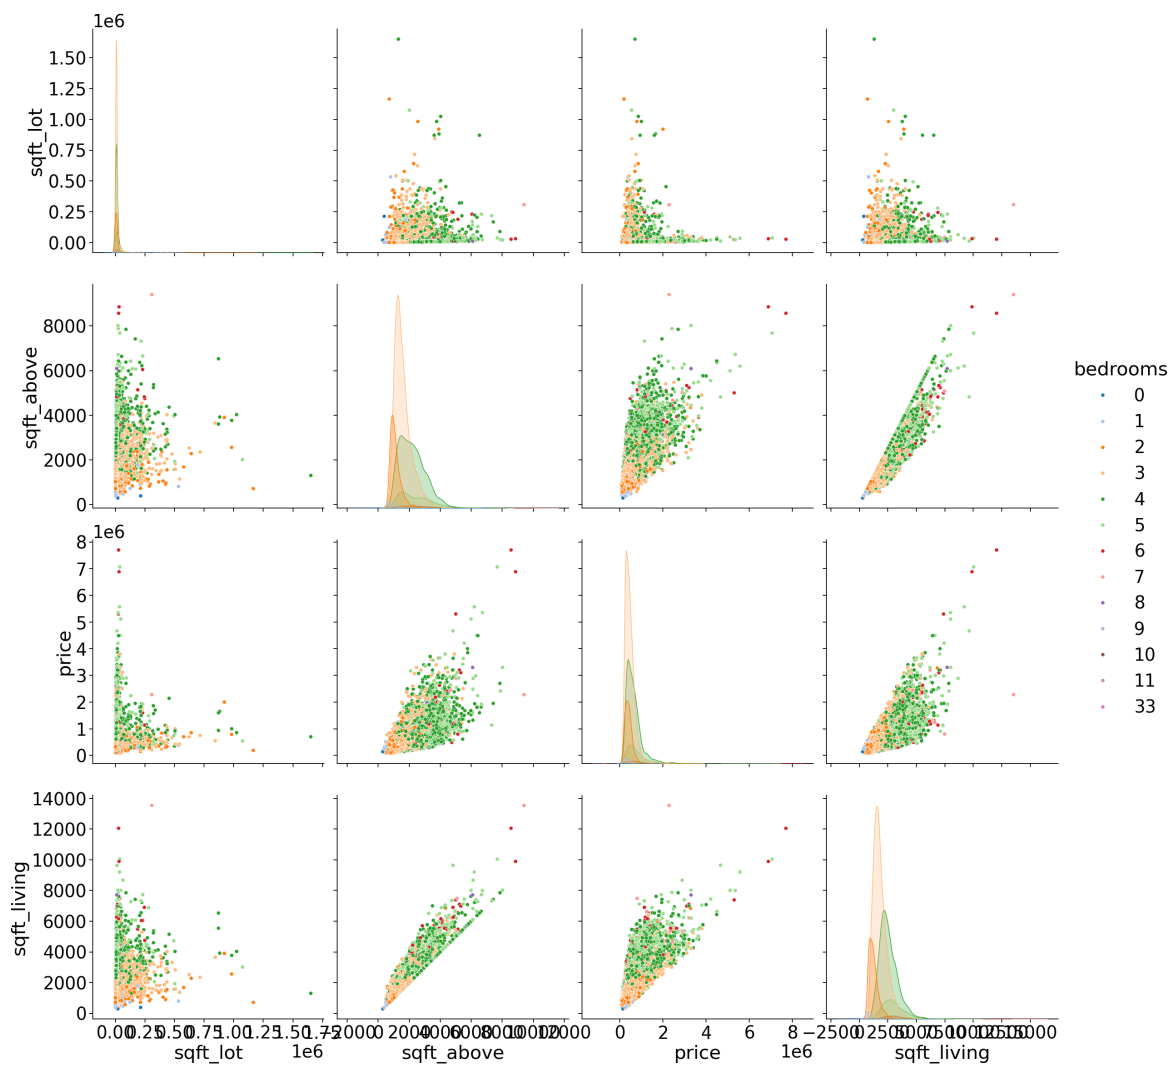
```
In [15]: dataset= dataset.drop(['id','date'], axis=1)
```


```
In [31]: with sns.plotting_context("notebook", font_scale=2.5):
          g= sns.pairplot(dataset[['sqft_lot','sqft_above','price','sqft_living','bedr
          g.set(xticklabels=[]);
          plt.show()
```


```


C:\Users\DELL\anaconda3\Lib\site-packages\seaborn\axisgrid.py:2100: UserWarning:
The `size` parameter has been renamed to `height`; please update your code.
  warnings.warn(msg, UserWarning)


```




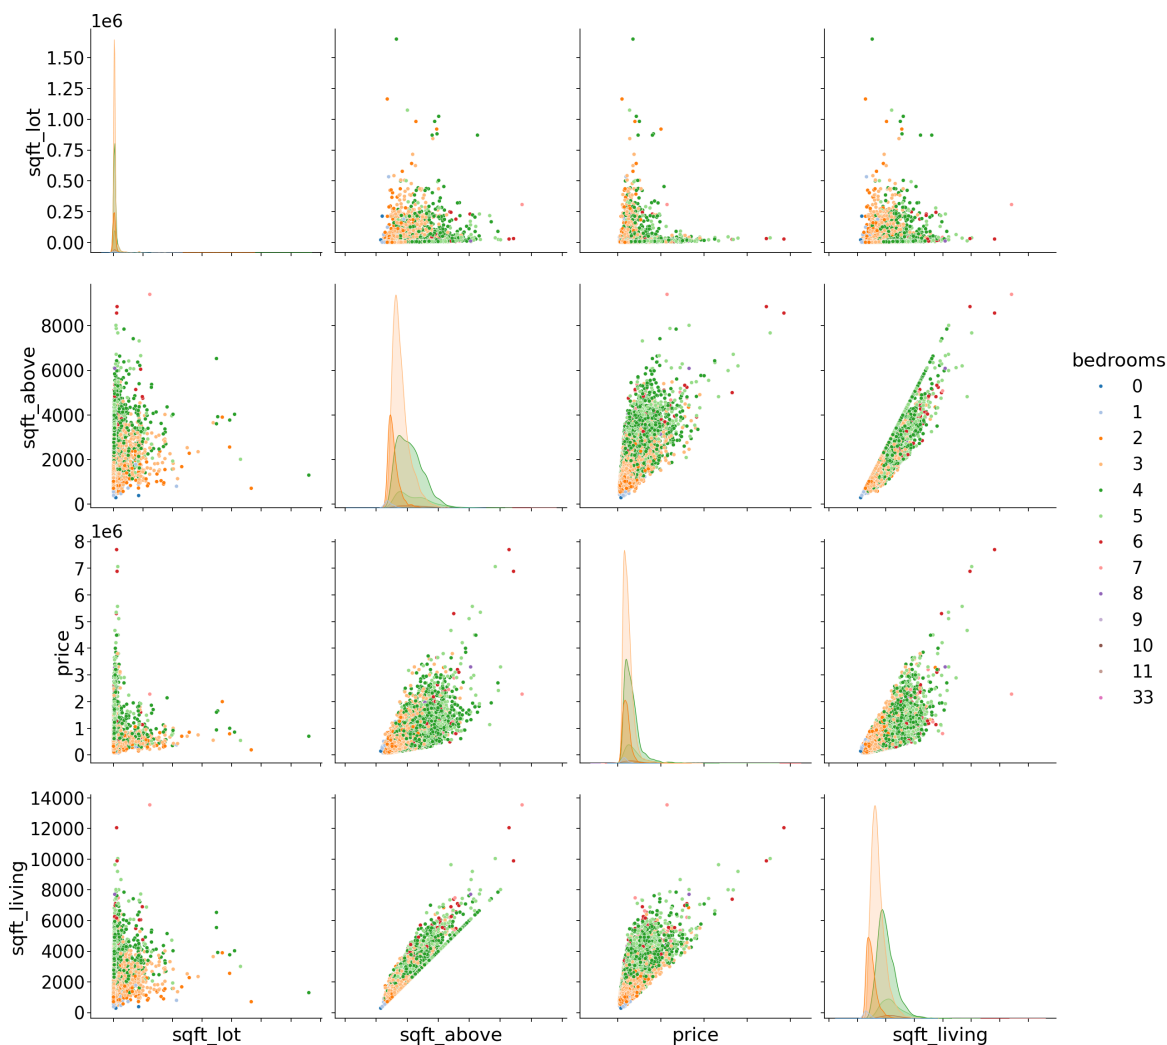












```
In [33]: X=dataset.iloc[:,1:].values
y=dataset.iloc[:,0].values

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size= 1/3, random_
```

```
In [37]: from sklearn.linear_model import LinearRegression
regressor= LinearRegression()
regressor.fit(X_train,y_train)
```

```
Out[37]: LinearRegression
LinearRegression()
```

```
In [39]: y_pred = regressor.predict(X_test)
```

```
In [41]: y_pred
```

```
Out[41]: array([ 386540.99847831, 1516969.01534114,  538662.72575254, ...,
 526000.75505741,  313924.63663316,  400525.6731456 ])
```

```
In [45]: import statsmodels.api as sm
def backwardElimination(x, SL):
    numVars = len(x[0])
    temp = np.zeros((21613,19)).astype(int)
    for i in range(0, numVars):
```

```

regressor_OLS = sm.OLS(y, x).fit()
maxVar = max(regressor_OLS.pvalues).astype(float)
adjR_before = regressor_OLS.rsquared_adj.astype(float)
if maxVar > SL:
    for j in range(0, numVars - 1):
        if (regressor_OLS.pvalues[j].astype(float) == maxVar):
            temp[:,j] = x[:, j]
            x = np.delete(x, j, 1)
            tmp_regressor = sm.OLS(y, x).fit()
            adjR_after = tmp_regressor.rsquared_adj.astype(float)
            if (adjR_before >= adjR_after):
                x_rollback = np.hstack((x, temp[:,[0,j]]))
                x_rollback = np.delete(x_rollback, j, 1)
                print (regressor_OLS.summary())
                return x_rollback
            else:
                continue
regressor_OLS.summary()
return x

```

SL = 0.05

X_opt = X[:, [0, 1, 2, 3, 4, 5,6,7,8,9,10,11,12,13,14,15,16,17]]

X_Modeled = backwardElimination(X_opt, SL)

OLS Regression Results

```

=====
=====
Dep. Variable:          y    R-squared (uncentered):
0.905
Model:                OLS    Adj. R-squared (uncentered):
0.905
Method:                Least Squares    F-statistic:                1.2
11e+04
Date:                  Wed, 20 Aug 2025    Prob (F-statistic):
0.00
Time:                  15:20:47    Log-Likelihood:                -2.94
61e+05
No. Observations:      21613    AIC:                5.8
92e+05
Df Residuals:          21596    BIC:                5.8
94e+05
Df Model:              17
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
x1	-3.551e+04	1888.716	-18.802	0.000	-3.92e+04	-3.18e+04
x2	4.105e+04	3253.759	12.618	0.000	3.47e+04	4.74e+04
x3	110.2642	2.268	48.607	0.000	105.818	114.711
x4	0.1334	0.048	2.786	0.005	0.040	0.227
x5	5261.5471	3541.347	1.486	0.137	-1679.755	1.22e+04
x6	5.833e+05	1.74e+04	33.598	0.000	5.49e+05	6.17e+05
x7	5.236e+04	2128.298	24.600	0.000	4.82e+04	5.65e+04
x8	2.721e+04	2323.818	11.709	0.000	2.27e+04	3.18e+04
x9	9.548e+04	2145.492	44.503	0.000	9.13e+04	9.97e+04
x10	71.3928	2.238	31.902	0.000	67.006	75.779
x11	38.8714	2.624	14.813	0.000	33.728	44.015
x12	-2561.7953	68.006	-37.670	0.000	-2695.092	-2428.498
x13	20.4187	3.646	5.600	0.000	13.272	27.566
x14	-519.0756	17.826	-29.119	0.000	-554.016	-484.136
x15	6.022e+05	1.07e+04	56.106	0.000	5.81e+05	6.23e+05
x16	-2.179e+05	1.31e+04	-16.683	0.000	-2.44e+05	-1.92e+05
x17	23.0994	3.392	6.811	0.000	16.452	29.747
x18	-0.3761	0.073	-5.137	0.000	-0.520	-0.233

```

=====
=====
Omnibus:                18403.146    Durbin-Watson:                1.991
Prob(Omnibus):          0.000    Jarque-Bera (JB):            1873534.498
Skew:                   3.572    Prob(JB):                    0.00
Kurtosis:               48.049    Cond. No.                    4.80e+17
=====
=====

```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The smallest eigenvalue is 9.48e-22. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In []: